

# A brief introduction to DEI IBM POWER clusters

## Login to the server:

```
ssh username@power7a.dei.unipd.it
```

Maybe the system will ask you to change your password the first time you log in, otherwise use the `passwd` command to change it after the first usage.

## Basic usage:

Using the system might be tricky if you haven't ever used a batch system. You cannot have an interactive access to the computing resources: instead you have to request them through a 'job file' which goes along with your executable.

This job file have a particular syntax and must be *submitted* to the system scheduler. In the IBM world the scheduler is called LoadLeveler.

## LoadLeveler micro-primer:

Create your job file (simulation.job):

```
#---CUT HERE (simulation.job)---
#!/bin/bash
#@ job_type = parallel
#@ initialdir = /home/username/simulations
#@ input = /dev/null
#@ output = /home/username/simulations/my_output.txt
#@ error = /home/username/simulations/my_error.txt
#@ notification = error
#@ class = short
#@ blocking = UNLIMITED
#@ total_tasks = 20
#@ queue

./mysimulation -my_parameters
#---ENDS HERE---
```

## Dissection of simulation.job:

```
#!/bin/bash
#@ job_type = parallel
```

The first line is the 'command interpreter': the lines of your job file beginning with '#' are treated as comments (but those starting with `#@` are intercepted by LoadLeveler). Every other line is executed by the interpreter (bash in this case) as if it was typed on the command line.

job\_type could be 'parallel' or 'serial'.

```
#@ initialdir = /home/username/simulations
```

This is your 'base' directory containing at least *simulation.job* and the executable and where you would normally put your outputs. If the job file is in the same directory of all the other files you can also write

```
initialdir = .    (the period meaning 'my current directory').
```

```
#@ input = /dev/null
```

Normally it's easier to specify the input on the command line (see below) so you're stating here that you don't need any input.

```
#@ output = /home/username/simulations/my_output.txt
```

This is one of many ways to get your output: if your simulation writes on its own files it may be possible that you don't need this either. Some error messages coming from the system might go here also.

---

Tip: if you specified `initialdir = .` you can change the line in simply

```
#@ output = my_output.txt
```

```
#@ output=/home/username/simulations/my_error.txt
```

This is where errors you normally see on the screen go.

---

Tip: if you specified `initialdir = .` you can change the line in simply

```
#@ error = my_error.txt
```

```
#@ notification = error
```

The system will mail you in case something goes wrong.

```
#@ class = short
```

This is where LoadLeveler comes in: the 'class' directive tells the scheduler that your job will run for up to 2 minutes. These limits (2 minutes, maximum number of processors...) are hardcoded on the definition of the class. There are other classes you can use: you can list them all with the 'llclass' command (try it at least once!).

```
#@ blocking = UNLIMITED
#@ total_tasks = 20
```

You can safely ignore the meaning of the 'blocking' directive and use it as is.

The 'total\_tasks' is where you ask the system to reserve the CPUs for your

MPI/parallel programs.

```
#@ queue
```

Mandatory: it ends the block of requirements.

```
./mysimulation -my_parameters
```

This is the command line you would normally type at the console. The path is relative to `initialdir`.

## Submitting your job:

At the console, type

```
$ llsubmit simulation.job
```

The system will answer something like:

```
llsubmit: The job "power7a.dei.unipd.it.28" has been submitted.
```

The number '28' is your '*job id*': the scheduler will queue your job and execute it as soon as possible. The global execution queue is listed by using the '`llq`' command. As an example:

```
$ llq
Id                        Owner      Submitted   ST PRI Class      Running On
-----
power7a.29.0             paolo      12/20 12:20 R   50  verylong      power7l

1 job step(s) in queue, 0 waiting, 0 pending, 1 running, 0 held, 0 preempted
```

Above you see that my job is running since the `ST` (state) column says 'R' and that I'm using the 'verylong' class.

If something goes wrong it may be possible that your remains stuck in the 'I' (idle) state: in that case you can use

```
llq -l -s your_job_id
```

to ask the scheduler what's happening.

Cancel any job you don't want to run anymore using

```
llcancel your_job_id
```

## Getting further documentation:

- `man command_i_know_nothing_of`

- IBM Cluster information center:

<http://publib.boulder.ibm.com/infocenter/clresctr/vxrx/index.jsp>

look for Tivoli Workload Scheduler LoadLeveler pages.

## Improving your job file:

Frequently you will do many repetitions of your simulations. Using the 'basic' version of *simulation.job* ends up in overwriting your output and error files on every run: can we do better? The answer is yes and the solution consists in 'numbering' automatically your files. You can use for this purpose your *jobid* or the current date and time:

- using *jobid*

```
#@ output = simulation_${jobid}.out.txt
#@ error = simulation_${jobid}.err.txt
```

- using the current date and time:

```
#!/bin/bash
...
... (other LoadLeveler directives)
...
now = `date +%Y%m%d-%H%M%S` (notice that ` are backticks characters)
#@ output = simulation_${now}.out.txt
#@ error = simulation_${now}.err.txt
```

## Interactive execution:

For a quick trial or for debugging purposes, you may want to launch your program like you would do with a serial one in an interactive session. This can be done by using a *host.list* file: this is a simple text file where you name the hosts (one per line) where you want your program to be launched. Since our Power cluster consist of only one (big) host, your *host.list* will be made of N or more 'power7a' equal lines I, N being the number of tasks you need. Of course any number  $M > N$  of lines will do.

Example: you want to run myprogram on 4 tasks on 'power7a' cluster.

- contents of *host.list*:

```
power7a
power7a
power7a
power7a
```

- from the shell:

```
$ ./myprogram -procs 4
```

Remember to put *host.list* in the same directory of *myprogram*